



Argo Workflows 101: Fundamentals

10am-noon PST

22nd Sep 2020

[Recording](#)



Pre-requisites

1. Install Kubernetes locally (we recommend use Docker on Desktop + K3D as it supports RBAC):
2. `kubectl`
3. [Add yourself to the sign-in sheet](#)

```
brew install k3d
```

```
k3d create ;# or `k3d cluster create` for newer versions
```

```
export KUBECONFIG="$(k3d get-kubeconfig --name='k3s-default')" ;# `or ? for new  
version?
```

```
kubectl cluster-info
```



Install Argo Workflows

```
kubectl create ns argo
```

```
kubectl -n argo apply -f https://raw.githubusercontent.com/argoproj/argo/master/manifests/quick-start-postgres.yaml
```

```
kubectl -n argo patch cm workflow-controller-configmap -p '{"data": {"containerRuntimeExecutor": "pns"}}' ;# needed for K3S
```

```
kubectl -n argo get pods --watch ;# takes maybe 2m for all pods to be ready
```

```
kubectl -n argo port-forward svc/argo-server 2746:2746
```

```
open http://localhost:2746
```

```
brew install argo
```

```
argo version
```

Fundamentals

Alex

Hands-On:

Using the user interface - submit a workflow that prints “Hi Argo Workshop!”



Hello Argo!

```
argo submit -n argo https://raw.githubusercontent.com/argoproj/argo/master/examples/hello-world.yaml
```

```
argo list -n argo
```

```
argo get -n argo ...
```

```
argo logs -n argo ...
```

Hands-On:

Using the CLI - submit a workflow and wait for it to finish.

Hint: https://argoproj.github.io/argo/cli/argo_submit/



Workflow Service Account

A short detour on running workflows with different service accounts.

```
kubectl create serviceaccount me
```

```
kubectl create rolebinding me --serviceaccount=argo:me --role=workflow-role
```


Hands-On:

Submit a workflow that uses a service account.

Anatomy of a Workflow

Simon



Anatomy of a workflow

- Steps
- DAGs
- Containers
- Scripts
- Resources suspend
- Arguments, Inputs, and Outputs
- Artifacts
- Exit handler



Templates

Templates are how we define the work to be done and call other templates to do the work.

All templates are defined under the `templates` field of a Workflow.

Templates act like functions/methods (we will see soon).

Several different templates, but two main kinds: those that define work and those that call on other templates (we will see soon).

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *define* work in methods

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *define* work in methods, give them names

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *define* work in methods, give them names, and specify their inputs and outputs


```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We call work in code blocks

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *call* work in code blocks by naming our desired definitions

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *call* work in code blocks by naming our desired definitions, passing in and receiving live arguments

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = addFour(2);  
        if (result > 5) {  
            sayHello();  
        }  
    }  
    public int addFour(int a) {  
        return a + 4;  
    }  
    public void sayHello() {  
        System.out.println("Hello Intuit!");  
    }  
}
```

We *call* work in code blocks by naming our desired definitions, passing in and receiving live arguments, and do some execution control

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 ) )"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

We *define* work in methods

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}}" > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
      template: addFour
      arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
      template: sayHello
      when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *define* work in methods, give them names

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *define* work in methods, give them names, and specify their inputs and outputs


```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *call* work in code blocks

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello(string toWho) {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *call* work in code blocks by naming our desired definitions

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *call* work in code blocks by naming our desired definitions, passing in and receiving live arguments

```

class ArgoDemo {
    public static void main(String[] args) {
        int result = addFour(2);
        if (result > 5) {
            sayHello();
        }
    }

    public int addFour(int a) {
        return a + 4;
    }

    public void sayHello() {
        System.out.println("Hello Intuit!");
    }
}

```

```

- name: main
  steps:
    - - name: addFour
        template: addFour
        arguments: {parameters: [{name: "a", value: "2"}]}
    - - name: sayHello
        template: sayHello
        when: "{{steps.addFour.outputs.result}} > 5"
- name: addFour
  inputs: {parameters: [{name: "a"}]}
  container:
    image: alpine:latest
    command: [sh, -c]
    args: ["echo $(( {{inputs.parameters.a}} + 4 )")"]
- name: sayHello
  container:
    image: alpine:latest
    command: [sh, -c]
    args: [echo "Hello Intuit!"]

```

We *call* work in code blocks by naming our desired definitions, passing in and receiving live arguments, and do some execution control



Templates

Templates are how we define the work to be done and call other templates to do the work.

Two kinds of templates:

Definition templates:

- Container (from the last example)
- Script
- Resource
- Suspend

Execution templates:

- Steps (from the last example)
- DAG



Inputs and Output Parameters

Input and Output Parameters are how we move data across different steps in a Workflow.

- Inputs are a *definition* of inputs `addFour(int a)`
- Outputs are a *definition* of outputs `int addFour(...)`
- Arguments are *live* arguments `addFour(2)`



Hands On

1. Grab our working example from:
<http://in/argo-hw> or <https://bit.ly/2ZyiPmO>
2. Modify it as defined in the new Java code (changes highlighted)
3. Challenge: Can you use a DAG template instead of Steps? (Docs and examples in the Argo Repo)

```
class ArgoDemo {  
    public static void main(String[] args) {  
        int result = add(2, 5);  
        int finalResult = add(result, 10);  
        if (finalResult > 5) {  
            sayHello(finalResult);  
        }  
    }  
    public int add(int a, int b) {  
        return a + b;  
    }  
    public void sayHello(int result) {  
        System.out.println("Result is: " + result);  
    }  
}
```

Artifacts

Bala



Input and Output Artifacts

- Artifacts can be a single file or directory
- Supported Repositories
 - S3, GCS, OSS, RAW, HDFS, Github, http
- Configure Repository in Argo
 - Controller level configuration in configmap
 - Workflow level configuration
 - ArtifactRepositoryRef
 - Inline configuration

```
s3:  
  bucket: my-bucket  
  endpoint: minio:9000  
  insecure: true  
  accessKeySecret:  
    name: my-minio-cred  
    key: accesskey  
  secretKeySecret:  
    name: my-minio-cred  
    key: secretkey
```



Inline artifact repository

```
- name: input-artifact-s3-example
  inputs:
    artifacts:
      - name: my-art
        path: /my-artifact
        s3:
          # Use the corresponding endpoint depending on your S3 provider:
          #   AWS: s3.amazonaws.com
          #   GCS: storage.googleapis.com
          #   Minio: my-minio-endpoint.default:9000
          endpoint: s3.amazonaws.com
          bucket: my-bucket-name
          key: path/in/bucket
          # accessKeySecret and secretKeySecret are secret selectors.
          # It references the k8s secret named 'my-s3-credentials'.
          # This secret is expected to have the keys 'accessKey'
          # and 'secretKey', containing the base64 encoded credentials
          # to the bucket.
          accessKeySecret:
            name: my-s3-credentials
            key: accessKey
          secretKeySecret:
            name: my-s3-credentials
            key: secretKey
```



Using ArtifactRepositoryRef:

- Configure multiple repositories in configmap and refer them in workflow using `artifactRepositoryRef:`

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: artifactory-repository-ref-
spec:
  entrypoint: main
  artifactRepositoryRef:
    key: minio
  templates:
    - name: main
      container:
        image: docker/whalesay:latest
        command: [sh, -c]
        args: ["cowsay hello world | tee /tmp/hello_world.txt"]
      outputs:
        artifacts:
          - name: hello_world
            path: /tmp/hello_world.txt
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: artifact-repositories
data:
  minio: |
    s3:
      bucket: my-bucket
      endpoint: minio:9000
      insecure: true
      accessKeySecret:
        name: my-minio-cred
        key: accesskey
      secretKeySecret:
        name: my-minio-cred
        key: secretkey
```

Hands-On:

Submit a workflow that use artifacts.



Output Artifacts

- Output will be pushed to configured Repository once main container is completed
- <https://gist.github.com/sarabala1979/a7190fe6f43996f2ee6a2d9877723aaa>

```
templates:
- name: whalesay
  container:
    image: docker/whalesay:latest
    command: [sh, -c]
    args: ["cowsay hello world | tee /tmp/hello_world.txt"]
  outputs:
    artifacts:
      - name: message
        path: /tmp/hello_world.txt
        s3:
          bucket: my-bucket
          endpoint: minio:9000
          insecure: true
          key: output/hello_world.txt
          accessKeySecret:
            name: my-minio-cred
            key: accesskey
          secretKeySecret:
            name: my-minio-cred
            key: secretkey
```



Input Artifacts

- Input artifact will be downloaded from configured Repository and saved in given path for main container to access it.
- <https://gist.github.com/sarabala1979/a2198b888a31269afb5fe08c0de3af1d/raw/947898cf8f68b544b7c456f0bc992627b1a2cca/input-artifact.yaml>

```
templates:
- name: input-artifact
  inputs:
    artifacts:
      - name: my-art
        path: my-artifact
        s3:
          bucket: my-bucket
          endpoint: minio:9000
          insecure: true
          key: output/hello_world.txt
          accessKeySecret:
            name: my-minio-cred
            key: accesskey
          secretKeySecret:
            name: my-minio-cred
            key: secretkey
  container:
    image: debian:latest
    command: [sh, -c]
    args: ["cat my-artifact"]
```



Passing Artifacts

<https://gist.github.com/sarabala1979/03dcc960371851dbc6f7c9ea23abd212>



Exit Handlers

- It is a destructor of workflow
- You can define Exit handler in Workflow level, Step/Dag level
- <https://gist.github.com/sarabala1979/f3cacd7fafd2378f577049182beeabaf/raw/126b4ba530bf1b473234dc36f9f569576c865294/exithandler.yaml>.

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: exit-handlers-
spec:
  entrypoint: intentional-fail
  onExit: exit-handler
```

```
- name: exit-handler
  container:
    image: docker/whalesay:latest
    command: [cowsay]
    args: ["Exit-Handler"]
```

```
protected void finalize()
{
  System.out.println("object is garbage collected ");
}
```


Hands-On:

Create workflow with an exit handler.

Workflow Templates

Bala



Change your workflow to a workflow template

- Instead of submitting whole workflow every time.
- You can store the workflow definition in cluster
- You can refer or submit definition multiple time.
- You just change kind from `workflow` to `WorkflowTemplate`
- Cli command for create workflow template
- `argo template create <>`

```
class WorkflowTemplate{
```

```
void template(int a, int b){  
    // do something  
}  
}
```

```
Class workflow {  
Public static void main(String[] args){  
    WorkflowTemplate wfTmpl = new  
WorkflowTempalte();  
    wfTmpl.template(1,2);  
wfTmpl.template(3,4);  
}
```

Hands-On:

Change your workflow to a template and submit the template.

<https://gist.githubusercontent.com/sarabala1979/5f5ecff0976908300d020c5cda9ad53c/raw/239b785009bbab87d161f3ce96bc6d1da392256d/workflowtemplate.yaml>

apiVersion: argoproj.io/v1alpha1

kind: WorkflowTemplate

metadata:

name: add-example-template

spec:

entrypoint: main

templates:

- name: main

steps:

- name: addFour

template: addFour

arguments: {parameters: [{name: "a", value: "2"}]}

- name: sayHello

template: sayHello

when: "{{steps.addFour.outputs.result}} > 5"

- name: addFour

inputs: {parameters: [{name: "a"}]}

container:

image: alpine:latest

command: [sh, -c]

args: ["echo \$(({{inputs.parameters.a}} + 4))"]

- name: sayHello

container:

image: alpine:latest

command: [sh, -c]

args: [echo "Hello Intuit!"]



Refer workflow template in workflow

<https://gist.githubusercontent.com/sarabala1979/9c188e5c425acf0d795eaf0b61e02729/raw/1cf5ba90c1d3e175190121c1c9f948cdd199f9f2/workflowtemplateRef.yaml>

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: add-example-
spec:
  workflowTemplateRef:
    name: add-example-template
```

<https://gist.githubusercontent.com/sarabala1979/9eb280fdab084ea92455b8f87c70146/raw/b538a8a27983ff53ed4c3e3b8b36752ede61a7d5/templateRef.yaml>

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: workflow-template-hello-world-
spec:
  entrypoint: whalesay
  templates:
    - name: whalesay
      steps:
        - name: call-whalesay-template
          templateRef:
            name: add-example-template
            template: addFour
          arguments:
            parameters:
              - name:
                value: "3"
```



Submit the workflow from workflow template

```
`argo submit --from workflowtemplate/add-example-template --watch`
```



Cluster workflow template

- ``argo cluster-template create <>``

```
apiVersion: argoproj.io/v1alpha1
kind: ClusterWorkflowTemplate
metadata:
  name: add-example-template
spec:
  entrypoint: main
  templates:
    - name: main
      steps:
        - - name: addFour
            template: addFour
            arguments: {parameters: [{name: "a", value: "2"}]}
          - - name: sayHello
              template: sayHello
              when: "{{steps.addFour.outputs.result}} > 5"
        - name: addFour
          inputs: {parameters: [{name: "a"}]}
          container:
            image: alpine:latest
            command: [sh, -c]
            args: ["echo $(( {{inputs.parameters.a}} + 4 ) )"]
        - name: sayHello
          container:
            image: alpine:latest
            command: [sh, -c]
            args: [echo "Hello Intuit!"]
```

Cron Workflows

Simon



Cron Workflows

CronWorkflows are Regular Workflows that run on a schedule. Converting one is easy.

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: hello-world-
spec:
  entrypoint: whalesay
  templates:
    - name: whalesay
      container:
        image: docker/whalesay:latest
        command: [cowsay]
        args: ["hello world"]
```

```
apiVersion: argoproj.io/v1alpha1
kind: CronWorkflow
metadata:
  generateName: hello-world-
spec:
  entrypoint: whalesay
  templates:
    - name: whalesay
      container:
        image: docker/whalesay:latest
        command: [cowsay]
        args: ["hello world"]
```

```
apiVersion: argoproj.io/v1alpha1
kind: CronWorkflow
metadata:
  generateName: hello-world-
spec:
  entrypoint: whalesay
  templates:
    - name: whalesay
      container:
        image: docker/whalesay:latest
        command: [cowsay]
        args: ["hello world"]
```

```
apiVersion: argoproj.io/v1alpha1
kind: CronWorkflow
metadata:
  generateName: hello-world-
spec:
  workflowSpec:
    entrypoint: whalesay
    templates:
      - name: whalesay
        container:
          image: docker/whalesay:latest
          command: [cowsay]
          args: ["hello world"]
```

```
apiVersion: argoproj.io/v1alpha1
kind: CronWorkflow
metadata:
  generateName: hello-world-
spec:
  schedule: "2 * * * *"
  timezone: "America/Los_Angeles"
  concurrencyPolicy: "Replace"
  workflowSpec:
    entrypoint: whalesay
    templates:
      - name: whalesay
        container:
          image: docker/whalesay:latest
          command: [cowsay]
          args: ["hello world"]
```

```
apiVersion: argoproj.io/v1alpha1
kind: CronWorkflow
metadata:
  generateName: hello-world-
spec:
  schedule: "2 * * * *"
  timezone: "America/Los_Angeles"
  concurrencyPolicy: "Replace"
  workflowSpec:
    entrypoint: whalesay
    templates:
      - name: whalesay
        container:
          image: docker/whalesay:latest
          command: [cowsay]
          args: ["hello world"]
```